### **FENIL GAJJAR**

# KUBERNETES DAILY TASKS

# KUBERNETES ARCHITECTURE

- COMPREHENSIVE GUIDE
- THEORY + PRACTICAL
- REAL TIME SCENARIO TASKS
- DAILY TASKS SERIES

Unlock the core of Kubernetes! Explore its architecture and discover how it powers modern cloud-native applications

**FOLLOW FOR MORE** 

**CONTACT US** 



# Unveiling Kubernetes

# **Architecture:**

The Blueprint of Modern

**Container Orchestration** 







# 🚀 Welcome Back to Our Kubernetes Journey! 🌍



First and foremost, a massive thank you to everyone who supported and engaged with my first post, Introduction to Kubernetes! Your feedback, comments, and enthusiasm have been truly motivating. This journey isn't just about learning Kubernetes—it's about building a strong DevOps mindset and mastering the tools that shape the modern cloud-native world.

Now, let's move forward to the **next crucial topic—Kubernetes Architecture!** 



Understanding how Kubernetes is structured is essential for effectively deploying, managing, and scaling containerized applications. In this post, we'll break down:

- The **key components** of Kubernetes and their roles
- How these components work together to ensure high availability, scalability, and automation
- Why this architecture is designed for **fault tolerance and resilience** in large-scale deployments

Whether you're a beginner or already familiar with Kubernetes, grasping its architecture will set the foundation for deeper concepts like networking, scheduling, and workload management.

and unlock the secrets of how Kubernetes orchestrates applications at scale. 🔥



# The Significance of Kubernetes Architecture in

# **Modern DevOps**

Kubernetes is more than just a container orchestration tool—its architecture is the backbone of scalability, automation, and resilience in modern cloud environments. Understanding this architecture is crucial because it directly impacts how organizations deploy, manage, and optimize their applications.

#### Why Does Kubernetes Architecture Matter?

#### Scalability & Performance Optimization

The distributed nature of Kubernetes allows seamless horizontal scaling of applications, ensuring workloads can handle increased demand dynamically.

#### High Availability & Fault Tolerance

Kubernetes follows a decentralized control plane and worker node model, ensuring that failures in one part of the system do not disrupt the entire application.

#### Automation & Efficient Resource Management

With components like the Kube-scheduler, Controller Manager, and Kubernetes API Server, workloads are automatically scheduled, monitored, and optimized, reducing manual intervention.

#### Consistency Across Multi-Cloud & Hybrid Environments

By abstracting infrastructure complexities, Kubernetes architecture enables applications to run **consistently across on-prem, cloud, and hybrid setups**, ensuring portability and flexibility.

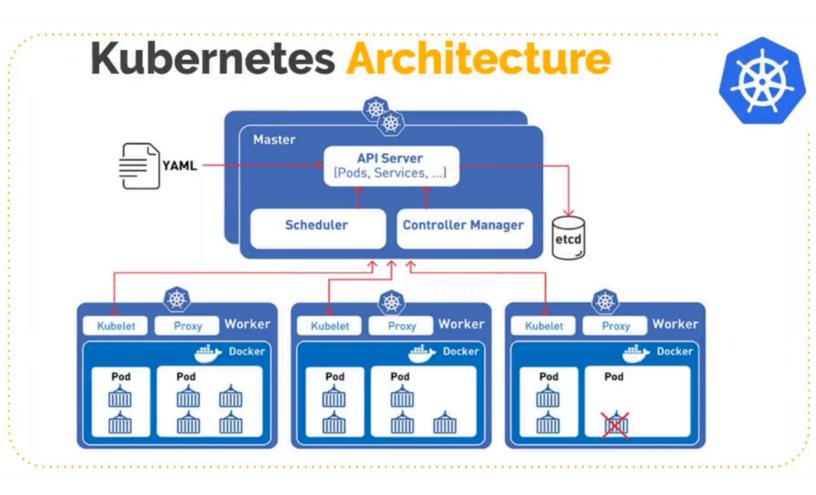
#### Security & Role-Based Access Control (RBAC)

A well-structured Kubernetes architecture ensures **secure cluster communication, fine-grained access controls, and robust security policies**, safeguarding applications against potential vulnerabilities.

# With the Core of Cloud-Native Innovation

A deep understanding of **Kubernetes architecture empowers DevOps teams** to design **scalable**, **resilient**, **and automated** deployment strategies, making it an **indispensable tool for cloud-native and enterprise-grade applications**.

# **T** Kubernetes Architecture: The Foundation of Scalable & Resilient Deployments





# Kubernetes Architecture Overview

Kubernetes follows a distributed architecture with two main categories of components:

- 1. Control Plane (Master Node) → Responsible for cluster management, scheduling, and monitoring.
- 2. Worker Nodes → Run the actual application workloads inside Pods (containers).

Both these components communicate over a **network layer** using Kubernetes APIs.

# 1 Control Plane (Master Node)

The Control Plane is responsible for making global decisions about the cluster. It maintains the desired state, schedules workloads, and ensures self-healing.

# **Components of the Control Plane**

#### 1. API Server (kube-apiserver)

- Acts as the entry point for all Kubernetes operations.
- Exposes the Kubernetes API, which users, controllers, and worker nodes interact with.
- Accepts REST API requests (kubect1 commands) and validates them.
- Updates the cluster state in ETCD and retrieves information when needed.

 Handles authentication and authorization using Role-Based Access Control (RBAC).

#### **Example:**

When you run kubectl apply -f deployment.yaml, the request goes to the API Server, which processes and updates the cluster state.

#### 2. ETCD (Key-Value Store)

- A highly available, distributed database that stores the entire Kubernetes cluster state.
- Stores configurations, secrets, and information about all Kubernetes objects (Pods, Deployments, Services, etc.).
- Runs on the Control Plane and must be backed up regularly to prevent data loss.

#### **Example:**

If a node crashes, Kubernetes retrieves the last known state from **ETCD** and recreates the workloads.

#### 3. Scheduler (kube-scheduler)

- Determines which worker node should run a new Pod.
- Monitors resource availability (CPU, memory) on worker nodes.
- Uses different scheduling policies like Affinity, Anti-Affinity, Taints, and
   Tolerations.
- Ensures workloads are evenly distributed across the cluster.

#### **Example:**

If a new Pod needs to be created, the Scheduler checks available nodes and assigns the Pod to the **best-suited node** based on CPU, memory, and constraints.

#### 4. Controller Manager (kube-controller-manager)

- Runs multiple controller processes to maintain the desired state of the cluster.
- Controllers watch the cluster state and take corrective actions if necessary.

#### **Types of Controllers**

- Node Controller: Detects node failures and replaces them.
- **Replication Controller:** Ensures the correct number of Pods are running.
- Service Account & Token Controller: Manages authentication tokens for service accounts.

#### **Example:**

If a node goes down, the Node Controller detects the failure and schedules new Pods on a healthy node.

#### 5. Cloud Controller Manager (Optional)

- Only required when running Kubernetes in a cloud environment (AWS, Azure, GCP).
- Integrates Kubernetes with cloud services like Load Balancers, Volumes, and Network Policies.

# **2** Worker Node Components

Worker Nodes are responsible for running application workloads inside **Pods**.

Each Worker Node consists of multiple components to ensure efficient container execution and networking.

## Components of a Worker Node

#### 1. Kubelet

- An agent running on every worker node that communicates with the Control Plane.
- Ensures that the containers in Pods are running as expected.
- Receives instructions from the API Server and executes them locally.
- Monitors the health of running Pods and restarts them if they fail.

#### **Example:**

If a Pod crashes, the Kubelet reports this failure to the API Server, which triggers a new Pod creation.

#### 2. Container Runtime

- Responsible for running containers inside Pods.
- Kubernetes supports multiple container runtimes:
  - Docker
  - containerd
  - o CRI-O

 The container runtime pulls container images, starts containers, and manages networking.

#### **Example:**

When a Pod is scheduled, the container runtime fetches the required container image from a container registry (Docker Hub, AWS ECR) and starts it.

#### 3. Kube-Proxy

- Manages network communication between different Pods, services, and external users.
- Implements iptables or IPVS rules to enable service discovery and load balancing.
- Ensures that traffic is correctly routed between Pods across different worker nodes.

#### **Example:**

When a user accesses an application, Kube-Proxy routes the request to the correct **Pod** behind a **Service**.

# **3** Kubernetes Networking

Kubernetes **networking model** ensures seamless communication between Pods, Services, and external users.

#### **Types of Networking in Kubernetes**

#### 1. Pod-to-Pod Communication

- Every Pod gets a unique IP address (from a network overlay like Flannel or Calico).
- Pods can communicate directly within a namespace.

#### 2. Service Networking

- A Service provides a stable IP address and load balancing for a set of Pods.
- Used when multiple replicas of a Pod are running behind a common endpoint.

#### 3. Ingress Networking

- Exposes services to external users using a reverse proxy (NGINX, Traefik, HAProxy).
- Allows host-based and path-based routing for multiple services.

#### **Example:**

- If frontend needs to communicate with backend, it does so through a Service.
- If an external user wants to access myapp.com, an Ingress Controller routes traffic to the correct service.

# 4 Kubernetes Objects

Kubernetes manages applications using different **objects**, which define how workloads should run.

#### **Important Kubernetes Objects**

- **Pod:** Smallest unit in Kubernetes, runs containers.
- Deployment: Manages the desired number of Pod replicas and allows rolling updates.
- **Service:** Provides networking and load balancing for Pods.
- ConfigMap & Secret: Stores environment variables and sensitive data.
- PersistentVolume & PersistentVolumeClaim: Provides persistent storage for stateful applications.

#### **Example:**

- A **Deployment** ensures that 3 replicas of an application are always running.
- A **Service** exposes the Deployment to users.

# 5 How Kubernetes Works - Step-by-Step

#### 1. User Deploys an Application

- o The user runs kubectl apply -f deployment.yaml.
- The request is sent to the API Server, which stores the configuration in ETCD.

#### 2. Scheduler Assigns Workloads

• The **Scheduler** selects the best worker node to run the Pod.

#### 3. Pods are Created and Managed

 The **Kubelet** on the worker node pulls the container image and starts the Pod.

#### 4. Networking is Set Up

• **Kube-Proxy** routes network traffic between Pods and Services.

#### 5. Scaling and Self-Healing

- o If demand increases, **Auto-scaling** creates more Pods.
- If a Pod crashes, **Kubelet** restarts it.

# **Kubernetes Architecture: A Storytelling Way**

#### The Story: Kubernetes Island

Imagine a bustling **island called Kubernetes**, where various **villages (applications)** live in harmony. The island is run by a highly organized government and a team of tireless workers who ensure everything functions smoothly. This island is magical, as it can expand or shrink to accommodate more villagers whenever needed.

Let's explore how Kubernetes Island works, its governance, and its citizens:

#### 1. The Central Government: Control Plane

The **Control Plane** is the central government of Kubernetes Island. It ensures that every village is running as planned, new villages are built properly, and resources are allocated efficiently.

#### **Key Members of the Government**

#### API Server (The Town Hall):

- The API Server is the town hall where villagers and external visitors (users) make their requests.
- If you want to build a new village or add more houses, you visit the town hall and submit your plans.
- Example: A farmer submits a request to grow crops (deploy a pod).
   The town hall processes and approves the request.

#### 2. Scheduler (The Builder):

- The **Scheduler** is like a skilled builder who decides where new villages (pods) will be constructed.
- It checks the available space and resources on the island and chooses the best spot for the new village.
- Example: If you request to build three houses (pods) for your workers, the builder ensures they're placed in areas with enough food, water, and land.

#### Controller Manager (The Overseer):

- The Controller Manager is the overseer who ensures that every village matches its blueprint.
- If a house collapses in a village, the overseer ensures it is rebuilt immediately.
- Example: If you ask for 3 houses in a village and one gets destroyed,
   the overseer makes sure it is replaced.

#### 4. etcd (The Record Keeper):

- **etcd** is the magical record keeper who remembers everything about the island—every village, house, and road.
- If the government forgets something, it checks with the record keeper.
- Example: If someone asks, "How many villages exist on the island?"
   etcd provides the exact number.

#### 2. The Tireless Workers: Worker Nodes

The **Worker Nodes** are the hardworking citizens of Kubernetes Island. These nodes perform all the actual tasks to keep the villages thriving.

#### **Key Components of a Worker Node**

#### 1. Kubelet (The Manager):

- The **Kubelet** is the manager on each worker node. It ensures the workers (containers) do their jobs properly and reports back to the government.
- Example: If a farmer is supposed to plant crops (run a container), the manager ensures the farmer does the task as expected.

#### 2. Kube-proxy (The Postmaster):

- The **Kube-proxy** is like the island's postmaster. It manages communication between villages (pods) and ensures messages are delivered.
- Example: If one village needs tools from another, the postmaster routes the request to the right place.

#### 3. Container Runtime (The Workers):

The Container Runtime is the workforce that does the actual job.
 These are the farmers, carpenters, and fishermen (containers)
 working hard to sustain the village.

 Example: The farmer plants crops, the carpenter builds houses, and the fisherman catches fish—these are the containers running inside pods.

#### 3. The Island's Roads: Networking

The villages are connected by an intricate network of roads. These roads allow villagers to visit each other and share resources. Kubernetes uses a **flat networking model**, ensuring all villages can communicate seamlessly.

- Villages to Villages: Pods communicate directly with each other using their addresses.
- Villages to Visitors: Services act as a gateway, allowing external visitors to interact with villages.

Example: A village hosting a marketplace (web app) sets up a road (service) to allow tourists (users) to visit and buy goods.

#### 4. The Magic: Self-Healing and Scalability

Kubernetes Island is magical because it can **heal itself** and **grow or shrink** as needed.

#### • Self-Healing:

- If a storm destroys a house (pod), the overseer (Controller Manager)
   rebuilds it.
- Example: If the database pod crashes, Kubernetes automatically creates a new one.

#### • Scalability:

- During a festival, the island needs more workers. Kubernetes adds new worker nodes to handle the increased workload.
- Example: During Black Friday, a web application scales up to handle millions of visitors and scales down afterward to save resources.

#### 5. The Builders' Blueprint: Declarative Approach

On Kubernetes Island, everything operates using blueprints (YAML files). Instead of micromanaging, you simply provide a plan, and the government ensures it is followed.

#### Example:

- You submit a plan stating: "I want 3 houses in Village A."
- Kubernetes takes care of building and maintaining those 3 houses.

Now, let's dive deeper into each component and explore how they work together to power Kubernetes!

#### **Kube-API Server in Kubernetes**

#### What is Kube-API Server?

The **kube-apiserver** is the **core component** of the Kubernetes control plane. It acts as a **gateway for all communication** in the cluster and is responsible for **handling REST API requests**, **validating them**, and **coordinating the cluster state**.

- It is the only component that directly interacts with the etcd database, which stores the cluster state.
- All internal and external Kubernetes operations go through the API server.

#### **How Kube-API Server Works?**

- Receives Requests Users (kubectl, dashboard, or other tools) and internal components (controller manager, scheduler, etc.) send API requests.
- 2. **Authentication & Authorization** Checks if the request is from a valid user and whether they have permission.
- 3. **Validation & Admission Control** Ensures the request follows Kubernetes rules before applying changes.
- 4. **Updates etcd** If valid, the API server updates etcd with the new cluster state.
- Sends Responses Returns the updated state to the requester and informs other control plane components.

## **Key Responsibilities of Kube-API Server**

- Handles All Kubernetes API Requests From kubectl, controllers, web dashboards, and external tools.
- Authenticates & Authorizes Users Ensures security through authentication (tokens, certificates, etc.) and authorization (RBAC, ABAC, etc.).
- Validates Objects Before Storing Ensures YAML manifests and configurations are correct before applying.
- Acts as a Frontend for etcd The only component that reads and writes to the etcd database.
- Communicates with Control Plane Components Passes information to the scheduler, controller manager, and other components.

## **How Kube-API Server Interacts with Other Components?**

#### 1. With etcd (Database)

- Reads cluster state (e.g., checking if a Pod exists).
- Writes updates when new resources are created or changed.

#### 2. With Kube-Scheduler

- The API server notifies the scheduler when a new Pod is created but has no assigned node.
- The scheduler decides the best node and updates the API server.

#### 3. With Controller Manager

- Controllers (e.g., ReplicaSet, Deployment) watch the API server for changes.
- If a Pod crashes, the controller requests the API server to create a replacement.

#### 4. With Kubelet (Node Agent)

- The API server **sends Pod definitions** to kubelet on worker nodes.
- Kubelet **reports back** on the Pod's health and status.

#### **Kube-API Server Architecture & Internals**

#### 1. Request Handling Pipeline

- 1. **Authentication** Verifies user identity (certificates, tokens, etc.).
- 2. Authorization Checks if the user has permission (RBAC, ABAC, etc.).
- 3. **Admission Controllers** Validate and mutate requests before storing them.
- 4. Object Validation & Schema Checks Ensures objects are well-formed.
- 5. **Persistence in etcd** Stores the final state of the cluster.

#### 2. API Server Internals

- Runs as a container inside the control plane node.
- Written in Go, designed as a RESTful API server.
- Handles high traffic efficiently using caching and watch mechanisms.

# **Example: How Kube-API Server Processes a Request?**

**Scenario**: A user creates a new Pod using kubectl.

sh

CopyEdit

kubectl run mypod --image=nginx

- Step-by-step flow:
- 1 kubectl sends a REST API request to kube-apiserver.
- 2 API server authenticates & authorizes the request.
- 3 Admission controllers validate the request (e.g., checking resource limits).
- 4 API server writes the Pod definition to etcd.
- **Scheduler detects the new Pod** and assigns it to a node.
- 6 Kubelet on that node pulls the image and starts the Pod.

#### How to Interact with Kube-API Server?

1 Check API Server Status

kubectl get componentstatuses

2 Access API Directly (Without Kubectl)

curl -k https://localhost:6443/version

#### 3 List Available API Resources

kubectl api-resources

#### 4 Get Cluster State in JSON Format

kubectl get pods -o json

## **Security Features of Kube-API Server**

- TLS Encryption Secures all communication between components.
- RBAC (Role-Based Access Control) Defines user permissions.
- Webhook Authentication & Authorization Allows external authentication systems.
  - Audit Logging Tracks all API requests for security monitoring.

# High Availability of Kube-API Server

In production, the API server should be **highly available** by running multiple instances behind a **load balancer**.

- Single-Node Cluster (Not Recommended for Production)
  - One API server instance running on a control plane node.

## Multi-Master Cluster (HA Setup)

- Multiple API servers run on different control plane nodes.
- A **load balancer** distributes traffic between them.

# retcd in Kubernetes

#### What is etcd?

etcd is a **distributed key-value store** used by Kubernetes to store **all cluster data**, including:

- Cluster state (nodes, pods, deployments, etc.)
- Configuration data (ConfigMaps, Secrets)
- Networking details (Services, Endpoints)
- ✓ Authentication & authorization info (RBAC, users, roles)

It's a **highly available, strongly consistent** database based on the **Raft consensus** algorithm.

## Why is etcd Important in Kubernetes?

Kubernetes completely relies on etcd to function.

- If etcd goes down, the Kubernetes API server stops working (since it can't read/write data).
- If etcd loses data, the entire cluster state is lost.

#### How etcd Works in Kubernetes?

- **1** kube-apiserver is the only component that interacts with etcd.
- 2 Any change in Kubernetes (e.g., creating a Pod) is sent to the API server.
- 3 The API server writes the change to etcd.

4 Other components (kube-scheduler, controllers) read changes from etcd and act accordingly.

#### etcd Data Structure

etcd stores data in a simple **key-value format**. Example:

Key	Value
/registry/pods/default /my-app-pod	Pod details (YAML)
/registry/nodes/ip-10- 0-0-1	Node details
/registry/configmaps/m y-config	ConfigMap data

To view etcd data, you can use:

ETCDCTL\_API=3 etcdctl get / --prefix --keys-only

# High Availability (HA) Setup for etcd

In production, etcd is **deployed in a cluster** for **fault tolerance**.

- A minimum of 3 nodes is recommended (always use an odd number to avoid split-brain issues).
- etcd uses the **Raft algorithm** to maintain consistency.

#### **Example etcd Cluster Setup**

- 3 etcd nodes running on control plane nodes
- etcd config file (/etc/etcd/etcd.conf.yaml):

```
name: etcd-1
initial-cluster:
etcd-1=https://10.0.0.1:2380,etcd-2=https://10.0.0.2:2380,et
cd-3=https://10.0.0.3:2380
initial-cluster-state: new
listen-peer-urls: https://10.0.0.1:2380
listen-client-urls: https://10.0.0.1:2379
```

• If one etcd node **fails**, others can still **maintain cluster state**.

#### Best Practices for etcd in Kubernetes

Always take etcd backups

ETCDCTL\_API=3 etcdctl snapshot save /backup/etcd-backup.db

- Run etcd on separate dedicated nodes (if possible)
- ✓ Use TLS encryption for etcd communication
- Monitor etcd health using Prometheus & Grafana

# How to Backup & Restore etcd in Kubernetes

Taking regular etcd **backups** is crucial for disaster recovery. Here's how you can **backup and restore etcd** step by step.

# Step 1: Set Environment Variables

Set up environment variables for **etcdctl** (Run this on a control plane node):

export ETCDCTL\_API=3

export ETCDCTL\_ENDPOINTS="https://127.0.0.1:2379"

export ETCDCTL\_CACERT="/etc/kubernetes/pki/etcd/ca.crt"

export ETCDCTL\_CERT="/etc/kubernetes/pki/etcd/server.crt"

export ETCDCTL\_KEY="/etc/kubernetes/pki/etcd/server.key"

# ✓ Step 2: Take a Backup of etcd

Run this command to save a snapshot of etcd:

etcdctl snapshot save /backup/etcd-backup.db

Verify the backup file:

ls -lah /backup/etcd-backup.db

Check snapshot status:

etcdctl snapshot status /backup/etcd-backup.db

# **✓** Step 3: Restore etcd from Backup

If your etcd data is lost or corrupted, you can restore it from a backup.

#### 1 Stop the etcd Service

On all control plane nodes, stop etcd before restoring:

sudo systemctl stop etcd

#### 2 Restore the etcd Snapshot

Run this command to restore from backup:

etcdctl snapshot restore /backup/etcd-backup.db --data-dir
/var/lib/etcd-new

#### 3 Replace Old Data with Restored Data

Move the old etcd data out and replace it with the new data:

sudo mv /var/lib/etcd /var/lib/etcd-old
sudo mv /var/lib/etcd-new /var/lib/etcd

#### 4 Start etcd Service Again

Restart etcd to load the restored data:

sudo systemctl start etcd

# **✓** Step 4: Verify the etcd Cluster

Check if etcd is running properly:

etcdctl endpoint health

Check etcd members:

etcdctl member list

# ✓ Automating etcd Backups with a Cron Job

You can schedule automatic etcd backups with a cron job:

sudo crontab -e

Add this line to run a backup every night at 2 AM:

0 2 \* \* \* etcdctl snapshot save /backup/etcd-(date +%Y%m%d).db

# Kube-Scheduler in Kubernetes

## What is Kube-Scheduler?

Kube-Scheduler is a **control plane component** responsible for **assigning pods to worker nodes** in a Kubernetes cluster. It **does not create pods**, but **decides where they should run** based on available resources, constraints, and scheduling policies.

#### How Kube-Scheduler Works

- 1 Pod Creation → A pod without a node assignment (.spec.nodeName is empty) is detected.
- ☐ Filtering Nodes → It filters out nodes that don't meet pod requirements.
- **3** Scoring Nodes → It ranks the remaining nodes based on a scoring system.
- **4** Assigning a Node → The highest-ranked node is selected.
- **⑤** Binding the Pod → Kube-Scheduler informs the API server, which updates the pod's nodeName.

# Scheduling Process in Detail

1 Filtering (Predicates)

The scheduler first eliminates nodes that **don't meet the pod's requirements**:

- ▼ Taints & Tolerations Does the node allow the pod to run?
- **V** Node Selectors & Affinity Rules − Does the pod prefer a specific node?

- **Resource Availability** Does the node have enough CPU & memory?
- Pod Anti-Affinity Does the pod avoid certain nodes?
- **✓ Node Conditions** Is the node in a healthy state?

## 2|Scoring (Prioritization) 🜟

After filtering, the scheduler scores nodes based on different criteria:

- Least Requested Resources Prefers nodes with the most free resources.
- **✓ Node Affinity** Prefers nodes matching a pod's affinity rules.
- **✓ Spread Constraints** Distributes pods evenly across zones/nodes.
- ✓ Custom Plugins Users can add custom scheduling logic.

#### 3 Binding the Pod to a Node 🔗

After selecting the best node, Kube-Scheduler **updates the API server**, setting the pod's .spec.nodeName. The kubelet on that node then pulls the pod's container images and starts the pod.

# Example: How Kube-Scheduler Works in a Real Scenario

Let's say you create a pod:

apiVersion: v1

kind: Pod

metadata:

```
name: my-app
spec:
  containers:
    - name: my-container
     image: nginx
  nodeSelector:
     environment: production # Scheduler will look for nodes
with this label
```

- The scheduler checks available nodes with the environment=production label.
- If multiple nodes match, it scores them based on available CPU, memory,
   and affinity rules.
- The **best node is chosen**, and the pod is assigned to it.

#### Advanced Kube-Scheduler Features

# **✓** Node Affinity (Preferred Nodes)

Example: Schedule a pod on nodes with disktype=ssd:

```
affinity:
```

nodeAffinity:

```
nodeSelectorTerms:
      - matchExpressions:
        - key: disktype
           operator: In
           values:
           - ssd
V Taints & Tolerations (Controlling Pod Placement)
Taint a node:
kubectl taint nodes node1 key=value:NoSchedule
Pod toleration:
tolerations:
  - key: "key"
    operator: "Equal"
    value: "value"
    effect: "NoSchedule"
```

requiredDuringSchedulingIgnoredDuringExecution:

#### Custom Schedulers

You can run your own scheduler instead of Kube-Scheduler:

kubectl get pods --all-namespaces -o wide | grep scheduler

# High Availability for Kube-Scheduler

- Run multiple schedulers (in an HA control plane).
- Use leader election to avoid conflicts.
- Monitor scheduler health with Prometheus & logs.

# **Kube-Controller Manager in Kubernetes**

#### What is Kube-Controller Manager?

The **kube-controller-manager** is a core component of Kubernetes that **runs multiple controllers in a single process**. These controllers **watch the cluster state** and ensure that the desired state (from manifests) matches the actual state.

**Example:** If a Pod crashes, the **ReplicaSet controller** automatically creates a new Pod.

#### Key Responsibilities of Kube-Controller Manager

The kube-controller-manager **runs multiple built-in controllers**, each responsible for different cluster resources.

#### 1 Node Controller

- Monitors the health of worker nodes
- If a node becomes unresponsive, it marks the node as NotReady
- If a node remains unresponsive for a long time, it removes Pods from that
   node

#### 📌 Example:

If a worker node stops responding for **5 minutes**, the Node Controller removes the Pods running on that node.

#### 2 Replication Controller

- Ensures that the desired number of Pod replicas are always running
- Works with ReplicaSet, Deployment, DaemonSet, etc.

#### **#** Example:

If a ReplicaSet defines **3 replicas**, but one Pod crashes, the controller **automatically creates a new Pod**.

#### 3 Endpoint Controller

- Manages the **Endpoints** object, which tracks the IPs of Pods behind a Service
- Ensures Services correctly route traffic to healthy Pods

#### **#** Example:

If a Pod crashes, the Endpoint Controller **removes its IP from the Service** and updates the list of healthy endpoints.

#### 4 Service Account & Token Controller

- Creates default ServiceAccounts for every new Namespace
- Generates API tokens for Pods to communicate with the Kubernetes API

#### **P** Example:

Every new Namespace automatically gets a default ServiceAccount.

#### **5** Persistent Volume Controller

- Manages Persistent Volumes (PVs) and Persistent Volume Claims (PVCs)
- Binds PVs to PVCs and ensures storage is allocated correctly

#### **\*** Example:

If a user requests 10GB storage (PVC), the controller binds it to an available PV.

#### 6 Job Controller

- Manages Job and CronJob resources
- Ensures that Jobs run to completion

#### **#** Example:

If a **batch job** is supposed to run once, the Job Controller ensures it completes successfully.

#### How Kube-Controller Manager Works?

- 1. Watches cluster objects (Nodes, Pods, Services, etc.)
- 2. Detects differences between desired and actual state
- 3. Takes **corrective action** using built-in controllers

#### **P** Example Flow:

- A Deployment specifies 5 Pods
- One Pod crashes → The Replication Controller detects it
- A new Pod is automatically created to maintain 5 Pods

# High Availability (HA) for Kube-Controller Manager

Since it is a critical component, HA is recommended.

#### **Mathematical Mathematical How to enable HA?**

- Run multiple kube-controller-manager instances across control plane nodes
- Use **leader election** (only one active instance at a time)
- If the active one fails, another takes over

#### Check active leader:

kubectl get endpoints kube-controller-manager -n kube-system

# Deploying Kube-Controller Manager Manually (Self-Managed Cluster)

If you're running a **self-managed Kubernetes cluster**, you can manually set up kube-controller-manager.

#### 1 Install kube-controller-manager

Download and install Kubernetes binaries:

```
wget
https://dl.k8s.io/release/v1.28.0/bin/linux/amd64/kube-contr
oller-manager
chmod +x kube-controller-manager
mv kube-controller-manager /usr/local/bin/
2 Create a systemd Service for Kube-Controller Manager
Create a service file
/etc/systemd/system/kube-controller-manager.service:
[Unit]
Description=Kubernetes Controller Manager
After=network.target
[Service]
ExecStart=/usr/local/bin/kube-controller-manager \
  --bind-address=0.0.0.0 \
  --cluster-signing-cert-file=/etc/kubernetes/pki/ca.crt \
  --cluster-signing-key-file=/etc/kubernetes/pki/ca.key \
  --kubeconfig=/etc/kubernetes/controller-manager.kubeconfig
```

```
--leader-elect=true
```

Restart=always

[Install]

WantedBy=multi-user.target

Start the service:

systemctl daemon-reload
systemctl enable kube-controller-manager

systemctl start kube-controller-manager

Note: In a managed Kubernetes cluster (e.g., EKS, AKS, GKE, kops), this setup is handled automatically.

In Kubernetes, **each controller** (like Node Controller, ReplicaSet Controller, etc.) is **logically a separate function** that manages different resources.

However, instead of running each controller as a separate process, Kubernetes combines all controllers into a single program called the kube-controller-manager. This helps reduce complexity and resource usage.

#### **★** Simple Analogy:

Think of each controller as a different worker in a company. Instead of having

separate offices for each worker, they all work in **one office**(kube-controller-manager) to make communication and management easier.

#### Why is this done?

- Less overhead (running multiple processes would be inefficient)
- Easier to manage (all controllers run under one service)
- Uses leader election (only one instance actively makes decisions)

# List of Controllers in Kube-Controller Manager

The **kube-controller-manager** runs multiple controllers, each responsible for managing different resources in the Kubernetes cluster.

#### 1 Node Controller

- Monitors the health of worker nodes
- Marks unresponsive nodes as NotReady
- Deletes Pods from a failed node after a timeout

#### **#** Example:

If a node stops responding for **5 minutes**, the controller removes its Pods.

#### 2 Replication Controller

- Ensures the desired number of Pod replicas are always running
- Works with ReplicaSet, Deployment, DaemonSet

#### **#** Example:

If a ReplicaSet defines **3 Pods**, and one crashes, it creates a new Pod.

#### **3** Endpoints Controller

- Updates the **Endpoints** object to keep track of healthy Pod IPs
- Ensures Services route traffic to the correct Pods

#### **P** Example:

If a Pod behind a Service crashes, the controller updates the Endpoint list.

#### 4 Service Account & Token Controller

- Creates a default ServiceAccount for each Namespace
- Generates API tokens for Pods to authenticate with the API server

#### **#** Example:

When a new Namespace is created, a **default** ServiceAccount is automatically added.

#### **5** Persistent Volume (PV) Controller

- Manages Persistent Volume (PV) and Persistent Volume Claim (PVC)
   bindings
- Handles dynamic storage provisioning

#### **#** Example:

If a user requests 10GB storage, the controller binds a PV to the PVC.

#### 6 Job Controller

- Ensures that **Jobs** and **CronJobs** run to completion
- Restarts failed Jobs if necessary

# **P** Example:

If a batch job is scheduled to **run once**, the controller ensures it completes.

#### **7** Garbage Collector Controller

- Cleans up unused or orphaned Kubernetes objects
- Detects and removes dangling Pods, Deployments, or Services

#### **#** Example:

If a Deployment is deleted, its ReplicaSet and Pods are also removed.

#### 8 Namespace Controller

• Deletes all resources inside a Namespace when the Namespace is deleted

#### **P** Example:

If dev-namespace is deleted, all Pods, Services, and Secrets inside it are also removed.

#### 9 Horizontal Pod Autoscaler (HPA) Controller

- Adjusts the number of Pods based on CPU/memory usage
- Works with metrics-server

#### **\*** Example:

If CPU usage exceeds 80%, it scales Pods from 2 to 4.

#### 10 Certificate Signing Controller

- Manages certificate requests for TLS security
- Signs Kubernetes API server certificates

#### **#** Example:

If a new worker node joins the cluster, this controller helps in securing communication.

#### Cloud Controller Manager (for Cloud Providers)

- Manages Kubernetes integration with cloud providers
- Handles Load Balancers, Volume Attachments, and Node Health

#### **P** Example:

If using AWS, this controller automatically provisions an **ELB (Elastic Load Balancer)**.

# Cloud Controller Manager (CCM) in Kubernetes

# **★** What is Cloud Controller Manager (CCM)?

The **Cloud Controller Manager (CCM)** is a Kubernetes component that allows the cluster to **interact with the underlying cloud provider** (AWS, GCP, Azure, etc.). It is responsible for managing cloud-specific operations like:

- Creating Load Balancers
- Attaching Persistent Volumes
- Managing Nodes (VMs/Instances)
- ✓ Handling Networking (Routes & IPs)

# 1 Why is CCM Needed?

Earlier, cloud provider-specific code was embedded in the **kube-controller-manager**, which made Kubernetes heavily dependent on cloud platforms.

To separate cloud-specific logic, Kubernetes introduced CCM, which allows:

- ✓ Decoupling Kubernetes from cloud providers
- ✓ Easier Kubernetes updates without worrying about cloud changes
- ✓ Running Kubernetes on-premise without unnecessary cloud logic

# 2 Key Responsibilities of Cloud Controller Manager

# **CCM** runs multiple controllers to manage cloud interactions:

#### 1 Node Controller

#### ★ Manages cloud-based Nodes (VMs/Instances)

- Checks if a Node (VM) is still active in the cloud
- If a Node is deleted in the cloud, the controller removes it from Kubernetes

#### **Example:**

If an EC2 instance is manually terminated in AWS, CCM automatically removes it from the cluster.

#### 2 Route Controller

#### Manages routing between Kubernetes Nodes

- Works in cloud environments that don't support automatic networking (e.g., GCE, AWS, Azure)
- Ensures correct IP routing between Pods and Nodes

# **Example:**

On AWS, it configures **VPC route tables** so Pods on different Nodes can communicate.

#### **3** Service Controller

**★** Manages Load Balancers for Services

- Provisions cloud-based Load Balancers for Service type:
   LoadBalancer
- Updates Load Balancer with correct backend Pod IPs

#### Example:

If you expose a Service in AWS with type: LoadBalancer, CCM provisions an AWS Elastic Load Balancer (ELB) and links it to the Service.

#### 4 Volume Controller

#### Handles Persistent Volume (PV) provisioning

- Works with cloud-based storage (EBS, GCE Persistent Disk, Azure Disk)
- Attaches & detaches volumes to Kubernetes Nodes

#### **Example:**

If a user requests **50GB storage** via a PersistentVolumeClaim (PVC), CCM provisions an **AWS EBS Volume** and attaches it to the Pod.

# 3 How Cloud Controller Manager Works?

- 1 When a **new Node (VM)** is created, CCM communicates with the cloud provider to check if the Node exists and fetches its details.
- 2 If a Service with type: LoadBalancer is created, CCM provisions a cloud-based Load Balancer and assigns it a public IP.
- If a Persistent Volume is requested, CCM provisions and attaches a **cloud** storage volume to the correct Node.
- 4 If a Node is deleted in the cloud, CCM removes it from Kubernetes.

# Where Are Kubernetes Control Plane Components Stored & Run?

All control plane components in Kubernetes—kube-apiserver, etcd, kube-scheduler, kube-controller-manager, and cloud-controller-manager (if applicable)—are typically run as containers inside the control plane node.

#### **How Are Control Plane Components Deployed?**

- In Self-Managed Clusters (Kubeadm, KOPS, Bare Metal)
  - These components run as static pods managed by kubelet.

The YAML manifests for these components are stored in:

#### /etc/kubernetes/manifests/

 Kubelet monitors these files and ensures the components are always running.

#### 2 In Managed Kubernetes Clusters (EKS, AKS, GKE, etc.)

- Control plane components are fully managed by the cloud provider.
- You don't have direct access to them.
- They are deployed across multiple nodes for high availability (HA).

# Control Plane Components & Their Storage/Execution Details

#### 1 Kube-API Server

- **Function:** Handles all Kubernetes API requests.
  - Runs as a container inside the control plane node.
  - Stores API data in etcd (database).
  - Manifests are in:

/etc/kubernetes/manifests/kube-apiserver.yaml

Logs can be checked via:

journalctl -u kube-apiserver

#### 2 etcd (Key-Value Store for Kubernetes Data)

- Function: Stores all cluster data (Pods, Deployments, ConfigMaps, etc.).
  - Runs as a container in control plane.
  - Data is stored in:

/var/lib/etcd/

- Uses Raft consensus algorithm for HA.
- Backup command:

etcdctl snapshot save /backup/etcd-snapshot.db

Logs can be checked via:

journalctl -u etcd

#### 3 Kube-Scheduler

- Function: Assigns Pods to Nodes based on resources & policies.
  - Runs as a container in control plane.
  - Manifests are in:

/etc/kubernetes/manifests/kube-scheduler.yaml

Logs can be checked via:

journalctl -u kube-scheduler

#### 4 Kube-Controller Manager

- Function: Runs multiple controllers (Node, Deployment, ReplicaSet, etc.).
  - Runs as a container in control plane.
  - Manifests are in:

/etc/kubernetes/manifests/kube-controller-manager.yaml

Logs can be checked via:

journalctl -u kube-controller-manager

#### 5 Cloud Controller Manager (Optional, for Cloud Environments)

- **Function:** Manages cloud integrations (Load Balancers, Nodes, Storage).
  - Runs as a container in control plane.
  - Manifests (if self-managed) are in:

/etc/kubernetes/manifests/cloud-controller-manager.yaml

Managed Kubernetes (EKS, GKE, AKS) already includes it.

#### **Hidden but Important Facts We Should Know**

#### **All these components run as containers inside the control plane node.**

You can verify using:

crictl ps

 This will show running control plane components as containerized workloads.

#### 2 They use host networking.

 Unlike normal Pods, these containers run with hostNetwork: true so they can directly communicate with the system.

#### **3** Static Pods are managed by kubelet.

 If the kube-apiserver crashes, kubelet automatically restarts it using the YAML from /etc/kubernetes/manifests/.

#### 4 No kube-proxy on control plane nodes.

 kube-proxy (which manages network traffic between Nodes) usually runs only on worker nodes.

#### **5** Control plane components must be highly available.

 In HA setups, control plane nodes are spread across multiple machines for fault tolerance.

#### all control plane components run as static pods.

These include:

- kube-apiserver
- etcd
- **kube-scheduler**
- kube-controller-manager
- **cloud-controller-manager** (if applicable)

#### Who Restarts Them if They Fail?

**Kubelet** is responsible for restarting static pods.

- Static pods are not managed by the Kubernetes API (i.e., no Deployment or ReplicaSet controls them).
- Instead, kubelet watches the static pod manifests stored in /etc/kubernetes/manifests/.
- If a static pod crashes or is deleted, kubelet will automatically restart it
  using the manifest file.

#### **How Does Kubelet Manage Static Pods?**

- 1 Kubelet runs on the control plane node.
- 2 It monitors the YAML files inside:

/etc/kubernetes/manifests/

3 If any static pod (like kube-apiserver) crashes or is removed, <b>kubelet detects it and starts it again</b> .  4 Kubelet uses the <b>container runtime</b> (like <b>containerd or Docker</b> ) to restart the failed container.
How to Check If Control Plane Pods Are Running?
Run this command on the control plane node:
kubectl get pods -n kube-system
You'll see all control plane components running as pods in the kube-system namespace.
How to Verify Static Pod Configurations?
Check the static pod manifests stored in:
ls /etc/kubernetes/manifests/
You'll see files like:
kube-apiserver.yaml
etcd.yaml

```
kube-scheduler.yaml
kube-controller-manager.yaml
These files define the static pods.
If a Control Plane Pod Keeps Crashing, What Should You Do?
1 Check logs:
kubectl logs -n kube-system <pod-name>
2 Check kubelet logs for errors:
journalctl -u kubelet -f
3 Manually restart the kubelet service:
systemctl restart kubelet
4 If a static pod YAML file is missing or corrupted, restore it from a backup.
```

# **Worker Node Components in Kubernetes**

A **worker node** is responsible for running application workloads in a Kubernetes cluster. It contains the necessary components to execute pods and communicate with the control plane.

#### Main Components of a Worker Node

#### 1 Kubelet (Node Agent)

#### What it does?

- Communicates with the API server to receive pod definitions.
- o Ensures that the containers inside the pod are running as expected.
- o Restarts failed containers if needed.
- Manages static pods on the node.
- Key Responsibility: Ensures that the node runs assigned workloads correctly.

#### 2 Kube-Proxy (Networking Component)

#### What it does?

- Maintains network rules on the node for pod-to-pod and pod-to-service communication.
- Uses iptables or IPVS to enable load balancing for services.
- Key Responsibility: Manages network communication inside the cluster.

#### **3** Container Runtime (Docker, containerd, CRI-O, etc.)

- What it does?
  - Pulls container images from registries.
  - o Creates, starts, stops, and manages containers inside pods.
- Key Responsibility: Runs application containers inside pods.

#### 4 Pods (Application Workloads)

- What they do?
  - Each worker node runs one or more pods, which contain one or more containers.
  - These are the actual applications deployed in the cluster.
- **Key Responsibility:** Runs the application workloads inside containers.

# **Kubelet in Kubernetes**

#### What is Kubelet?

Kubelet is an **agent running on each worker node** in a Kubernetes cluster. It is responsible for **managing the lifecycle of pods and containers** assigned to that node.

#### **Key Responsibilities of Kubelet**

#### 1 Registers the Node with the Control Plane

- When a new worker node joins the cluster, Kubelet registers the node with the API server.
- The API server then recognizes the node as ready to run workloads.

#### 2 Communicates with the API Server

- Kubelet constantly watches for pod assignments from the API server.
- It receives pod specifications (YAML) and ensures that they run on the node.

#### 3 Manages Pods and Containers

- **Ensures pods are running** as per the pod definition.
- If a pod crashes, Kubelet automatically restarts it.
- If a pod is deleted from the API server, Kubelet **removes it from the node**.

#### 4 Interacts with the Container Runtime

- Kubelet does not run containers directly—it uses a container runtime (like Docker, containerd, or CRI-O).
- It sends requests to the runtime to:
  - V Pull images
  - Create and start containers
  - Stop or delete containers when no longer needed

#### 5 Monitors Node and Pod Health

- Kubelet collects resource usage (CPU, memory, disk, network) and reports it to the API server.
- It performs liveness and readiness probes to check if a container is healthy.
- If a container fails multiple times, Kubelet can restart the pod or notify the API server.

#### 6 Handles Static Pods (If Configured)

- Kubelet monitors /etc/kubernetes/manifests/ for static pod YAMLs.
- If it finds a new static pod definition, it **creates and runs the pod**.
- These pods are **not managed by the API server** but only by Kubelet.

#### **How Kubelet Works?**

- 1 Kubelet **registers the node** with the Kubernetes API server.
- 2 It watches the API server for pod assignments.
- 3 When a pod is assigned to the node, Kubelet **instructs the container runtime** to run the pod.
- 4 It continuously monitors the pod's health and restarts failed pods if necessary.
- 5 Kubelet reports the node's status (CPU, memory, disk, network usage) to the

API server.

6 If the node fails, the API server stops scheduling new pods to that node.

#### Where is Kubelet Installed?

- Runs on every worker node and control plane node (for static pods).
- Installed as a binary or as a systemd service (/usr/bin/kubelet).
- Can be configured via /var/lib/kubelet/config.yaml.

# **Example Kubelet Command (Starting Kubelet)**

On a worker node, Kubelet can be started with:

```
kubelet --kubeconfig=/etc/kubernetes/kubelet.conf \
    --config=/var/lib/kubelet/config.yaml \
    --container-runtime=remote \
```

--container-runtime-endpoint=unix://run/containerd/containe
rd.sock

#### **Breakdown:**

- ✓ --config=/var/lib/kubelet/config.yaml → Kubelet config file
- ✓ --container-runtime=remote → Uses an external container runtime
- V

--container-runtime-endpoint=unix://run/containerd/containe
rd.sock → Connects to containerd

#### **How to Check Kubelet Logs?**

To troubleshoot Kubelet issues, check logs using:

journalctl -u kubelet -f

To restart Kubelet:

systemctl restart kubelet

#### Why Does Kubelet Run on Control Plane Nodes?

Even though the control plane is mainly responsible for managing the cluster, it also needs to run certain workloads. This is why Kubelet is present on control plane nodes.

#### 1 Runs Static Pods (like control plane components)

- Control plane components (kube-apiserver, etcd, kube-scheduler, kube-controller-manager) are static pods.
- Kubelet on the control plane node monitors and ensures these static pods are always running.

Static pod manifests are stored in:

/etc/kubernetes/manifests/

#### 2 Allows Control Plane Nodes to Run Regular Pods (Optional)

- If a cluster is small, control plane nodes can also be used as worker nodes.
- In such cases, Kubelet allows scheduling of normal workloads (pods) on control plane nodes.

To allow workloads on a control plane node, you need to remove the taint:

```
kubectl taint nodes <control-plane-node>
node-role.kubernetes.io/control-plane-
```

#### 3 Health Monitoring and Communication

- Kubelet on the control plane nodes reports node health to the API server.
- If the control plane node has issues, Kubelet informs the API server, which can take necessary actions.

# **Kube-Proxy in Kubernetes**

#### What is Kube-Proxy?

Kube-Proxy is a **networking component** that runs on every node in a Kubernetes cluster. It is responsible for **routing network traffic** between services and pods, ensuring seamless communication inside and outside the cluster.

#### **Key Responsibilities of Kube-Proxy**

#### Manages Service Networking

- Kubernetes services provide a stable IP for accessing pods.
- Kube-Proxy routes traffic to the correct pod behind the service.

#### 2 Implements Load Balancing

- When multiple pods back a service, Kube-Proxy distributes incoming traffic among them.
- Ensures high availability and efficient load distribution.

#### 3 Handles Cluster IPs

- When a service is created, Kubernetes assigns a ClusterIP.
- Kube-Proxy ensures that any request to this IP gets forwarded to the right pod.

#### 4 Supports Different Network Modes

Kube-Proxy operates in different **modes** for handling traffic:

- iptables mode (Default) Uses Linux iptables to forward traffic efficiently.
- **✓ IPVS mode (Better Performance)** Uses Linux IP Virtual Server (IPVS) for load balancing.
- ✓ Userspace mode (Deprecated) Uses a userspace process to forward traffic.

#### **How Kube-Proxy Works?**

- A service is created in Kubernetes.
- 2 Kube-Proxy sets up networking rules for that service using iptables or IPVS.
- 3 When a request comes to the service, Kube-Proxy forwards it to an available pod.
- 4 If a pod fails, Kube-Proxy **reroutes traffic** to healthy pods.

#### **Example: How Kube-Proxy Routes Traffic**

You create a service for a pod:

apiVersion: v1

kind: Service

metadata:

name: my-service

spec:

selector:

```
app: my-app
```

ports:

- protocol: TCP

port: 80

targetPort: 8080

#### 1. Kube-Proxy sets up networking rules

 Requests to my-service:80 are forwarded to an available pod's port 8080.

#### Where is Kube-Proxy Installed?

- Runs on every node (both worker and control plane nodes).
- Installed as a **DaemonSet** in kube-system namespace.

You can check its running status using:

kubectl get pods -n kube-system | grep kube-proxy

#### **How to Restart Kube-Proxy?**

If Kube-Proxy crashes, restart it with:

systemctl restart kube-proxy

or

kubectl delete pod -n kube-system -l k8s-app=kube-proxy

#### Is Kube-Proxy Running as a DaemonSet?

Yes, **Kube-Proxy runs as a DaemonSet** in Kubernetes by default.

- It is deployed in the **kube-system namespace**.
- Since it is a DaemonSet, it runs one instance per node to handle networking for that node.

You can verify its deployment with:

kubectl get ds -n kube-system | grep kube-proxy
or

kubectl get pods -n kube-system -l k8s-app=kube-proxy

#### Why is Kube-Proxy Referred to as "Optional" in Kubernetes Docs?

Kube-Proxy is **not strictly required** because its functionality can be replaced by **CNI plugins** that handle service networking differently. Here's why:

#### 1 Some CNI Plugins Handle Networking Without Kube-Proxy

 Modern CNI plugins like Cilium, Calico, and Antrea have built-in mechanisms for handling service traffic. These CNIs use eBPF (Extended Berkeley Packet Filter), which is faster
 than iptables/IPVS and eliminates the need for Kube-Proxy.

#### 2 eBPF-Based Networking is Becoming More Common

- eBPF is an advanced technology that allows direct kernel-level packet filtering and forwarding, making Kube-Proxy unnecessary in some setups.
- For example, Cilium with eBPF completely removes the need for Kube-Proxy while offering better performance.

#### **3** Some Kubernetes Deployments Don't Use Services for Traffic Routing

 If a Kubernetes cluster only uses Ingress controllers, service meshes (Istio, Linkerd), or external load balancers, Kube-Proxy may not be needed.

#### Should You Disable Kube-Proxy?

- If using a traditional iptables/IPVS-based setup, Kube-Proxy is required.
- If using eBPF-based CNIs (like Cilium), Kube-Proxy can be disabled.
- In most standard Kubernetes clusters, Kube-Proxy is still used for service networking.

To Disable Kube-Proxy (if using eBPF CNIs like Cilium)

kubectl -n kube-system delete ds kube-proxy

Note: Only disable Kube-Proxy if your CNI plugin explicitly supports it!

Yes, Kube-Proxy runs on control plane nodes as well.

#### Why Does Kube-Proxy Run on Control Plane Nodes?

- If control plane nodes also act as worker nodes (i.e., they schedule workloads), then Kube-Proxy is required to manage networking for pods running on them.
- Even if the control plane nodes don't schedule workloads, they may still need Kube-Proxy to ensure proper communication between Kubernetes services (like kube-apiserver, etcd, etc.).

#### How to Check if Kube-Proxy is Running on Control Plane Nodes?

You can verify this with:

kubectl get pods -n kube-system -o wide | grep kube-proxy

This will list all the nodes running Kube-Proxy, including control plane nodes.

#### Can We Disable Kube-Proxy on Control Plane Nodes?

 If control plane nodes don't run workloads, you can disable Kube-Proxy on them to reduce resource usage.

To do this, **taint control plane nodes** so no pods (other than system components) run there:

kubectl taint nodes <control-plane-node-name>
node-role.kubernetes.io/control-plane=:NoSchedule

 However, it's generally safe to keep Kube-Proxy running on control plane nodes to avoid any unexpected network issues.

Kubelet runs only as a daemon service on every node in the Kubernetes cluster. It is not deployed as a Kubernetes pod like other components.

#### Why Does Kubelet Run as a Daemon Service?

#### 1. Directly Managed by the OS

- Kubelet runs as a systemd service (or an init service) on Linux-based nodes.
- It starts when the node boots and ensures that workloads (pods) are running.

#### 2. Critical for Node Functionality

- Since Kubelet is responsible for communicating with the API server,
   it must always be running for the node to function in the cluster.
- If Kubelet were a Kubernetes pod, it would depend on itself to run,
   creating a failure loop.

#### 3. Not a Kubernetes Pod

 Unlike other control plane components (which run as static pods on control plane nodes), Kubelet does not run inside Kubernetes itself.  Instead, it is installed manually or via tools like kubeadm, kops, or managed services.

#### How to Check if Kubelet is Running?

You can check the status of Kubelet with:

systemctl status kubelet

or

ps aux | grep kubelet

If Kubelet is not running, you can start it using:

systemctl start kubelet

#### Can Kubelet Be Run as a Pod?

X No, Kubelet **cannot be a pod** because it is responsible for running and managing all other pods. Running Kubelet as a pod would create a **circular dependency issue**.

However, Kubelet does **register itself as a node** in the Kubernetes cluster, but it remains an external daemon process.

Kubernetes is a **powerful container orchestration system** that automates **deployment, scaling, and management** of containerized applications. It follows a **master-worker architecture**, where the **Control Plane** manages the cluster, and **Worker Nodes** handle workloads.

#### Request Handling & State Management

When a user submits a request (like deploying an application), it first goes to the **Kube-API Server** , which acts as the central hub of the cluster. The request is **validated** and then stored in **etcd** , a distributed key-value store that keeps the entire cluster state.

#### Scheduling & Workload Execution

The **Kube-Scheduler** (a) selects the best worker node based on available resources, then the **Kubelet** (a) on that node pulls the required container images and ensures the pod is running. Meanwhile, **Kube-Proxy** (a) manages networking, enabling seamless communication between pods and services.

#### Cluster Monitoring & Self-Healing

Controllers within the **Kube-Controller-Manager** continuously monitor the cluster and enforce the desired state. If a pod crashes or a node fails , Kubernetes automatically reschedules workloads to a healthy node, ensuring high availability and fault tolerance.

#### High Availability & Reliability

For **high availability**, multiple control plane nodes **iii ?** ensure redundancy through **leader election mechanisms** and a **distributed etcd cluster .** Core

components like the API server, scheduler, and controllers run as **static pods**, allowing the **Kubelet** to restart them if they fail.

With its **self-healing capabilities**, **automated scaling**, and **declarative management**, Kubernetes provides a **robust, scalable, and resilient** platform for running modern cloud-native applications.

# 

As we conclude this deep dive into **Kubernetes architecture**, we have explored how its core components seamlessly work together to orchestrate, manage, and scale containerized applications. From understanding the **control plane** and its critical role in cluster management to the **worker nodes** executing workloads efficiently, we have uncovered the inner workings of Kubernetes in detail.

This knowledge is essential for anyone looking to build, manage, or optimize Kubernetes clusters, whether for **self-managed environments** or **cloud-managed services**. Understanding the **API server, scheduler, controllers, etcd, kubelet, and networking components** provides a solid foundation to confidently navigate the Kubernetes ecosystem.

#### ★ What's Next?

Kubernetes is vast, and learning is an ongoing process. **Mastery comes from consistent practice, hands-on experience, and staying updated with evolving best practices**. That's why I'll be sharing **daily Kubernetes tasks**, practical challenges, and insights to help you deepen your understanding and sharpen your skills.

- Follow me for structured, real-world Kubernetes learning.
- Expect daily tasks, best practices, troubleshooting tips, and in-depth explanations.
  - Engage, ask questions, and let's grow together as a DevOps community.

I truly appreciate your time and support throughout this journey. **Your feedback, questions, and discussions have been invaluable.** Thank you for being part of this learning experience!

Let's continue exploring Kubernetes, DevOps, and cloud technologies together.



Stay tuned for more! #KeepLearning #Kubernetes #DevOps #CloudComputing